The testing and training data is split using Train and Test Split.py and it splits 80% of the data for training, and 20% for testing, and it ensures that the target distribution of both data sets is roughly equivalent, both with eachother and with the original dataset.

The split of 80/20 was chosen so the majority of the data can be used to test the model, while leaving 20% left (~900 records) to test the model and see what improvements have been made over the baseline model.